RESEARCH REPORT · ROUND 2 · OCTOBER 16, 2026

Build One Narrow Agent, Not the Platform

An AI-agent product for an independent NYC restaurant: what to build, what it costs, what already exists, what failed, how not to fail, and the honest truth about AI search recommendations.

> Written in simple English. Every number has a year and a clickable source. Vendor claims are labeled. Things we could not confirm are marked "not confirmed".

Contents (click to jump)

1. What we already know (from round one)
2. What you could build
3. What it costs to run per restaurant
4. Solutions that already exist
5. What people built that failed and why
6. How to make sure owners like it and it does not fail
7. Getting recommended by AI search tools: the honest truth
8. What changed or got corrected since round one
9. What we still could not verify
10. Bottom line

Bottom line up front. You can build a working AI-agent product for one NYC restaurant today, and the parts are cheap: about **$20–25 a month** per restaurant for a lean design and about **$275 a month** with a phone-answering voice agent (2026 list prices, Anthropic, Retell, Twilio). But an all-in-one is the wrong shape. Toast and Square already bundle nine of the twelve problems for $0–$699 a month, and the three they do not cover (portion size, "sweethearting" as behavior, phone use) need cameras, which run into a public accuracy ceiling and New York law. The one product that is missing, has peer-reviewed proof of payback, needs no hardware, and matches what owners already do by hand is a **POS-data monitor for voids, comps, discounts, refunds and no-sale drawer opens** (POS = the cash-register software; a void = cancelling an item after it was rung up; a comp = giving an item free). Build that first on Square, under $100 a month, month-to-month, with a dollar pass line set before the pilot starts. Owners in 2025–2026 do not leave because a tool lacked features; they leave because add-on fees grew, cancellation was blocked, and support vanished after the sale (Capterra reviews, 2025–2026). And no AI search tool can be made to recommend you "#1": repeated identical prompts return the same brand list less than 1% of the time (SparkToro, 2026), nobody sells placement inside the answer, and Gamma is a slide maker, not a recommender.

1. What we already know (from round one)

Round one established five things. The problems are structural: median restaurant profit was **2.8% of sales for full-service in 2024** and 42% of operators were not profitable in 2025 (NRA 2026), so no manager hour is spare. The buyer is unstable: **26% of restaurants change hands in year one and 60% by year three** (Parsa et al., 2005, study). Toast (about 180,000 locations as of June 2026, SEC 8-K) and Square already cover most of the list, including labor agents from $60 a month and marketing agents from $699 (Toast, October 2026, NRN). Cameras fail on three fronts: the best camera-only portion estimate has **18.8% mass error** (Nutrition5k, 2021, arXiv); New York forbids deducting losses from wages (Labor Law §193, NY DOL) and gates monitoring behind written notice (§52-c); and a 2.8%-margin owner cannot pay for hardware. And the one intervention with peer-reviewed payback is POS-data theft monitoring: across 392 restaurants, theft fell about 22% and revenue rose about 7% (Pierce et al., 2015; WashU, 2013), and 56% of 83,000 servers committed detectable theft at least once (WashU, 2019). Round one's call: do not build the all-in-one; build the void-and-comp monitor on Square, under $100 a month.

2. What you could build

The realistic design is a set of small agents that read data the restaurant already produces and send the owner a short message. The rule is "rules first, model second": plain code computes ratios (voids as a share of sales, hours clocked versus scheduled), and the LLM (large language model, the AI that writes text) only explains the handful of anomalies. That keeps the AI bill to a few dollars a month.

Agent 1 — Comp-and-void monitor (employee theft, sweethearting, employees taking money). Data: every order, discount, refund, void and no-sale drawer open, with the employee login and the manager code that approved it. Output: a ranked digest two or three times a week per login and shift, with time-stamped evidence. Known gap: on Square, a voided check comes back as "CANCELED" with no items;

Square staff confirmed "voided items aren't currently available with the Orders API" (2023, still open January 2025, Square forum). So on Square you can count voids and analyze discounts and refunds, but not see what was voided. Test comps on a live account before promising them.

Agent 2 — Clock checker (time theft, scheduling). Data: timecards and schedules from Square's Labor API (API = a door that lets one program read another program's data) or from 7shifts, whose token any admin can create with no approval (7shifts developers). Output: missed punches, early clock-ins, overtime drift. Most "time theft" is process failure: 39% of payroll errors came from missed punches (2026, Connecteam). Square Plus already blocks early clock-ins for $49 a month (Square), so this is a feature, not a product.

Agent 3 — Phone and customer-data agent (customer data, missed calls). A voice agent on Retell, Vapi or Bland answers calls, takes reservations, and captures the caller's name and number with consent. Cost: about **$106–179 a month for 1,200 minutes** (2026 list prices, Retell). It competes with Yelp Host at $149 a month (2025, Business Wire) and Slang.ai at $399+ (vendor claim). Delivery apps will not give you the data: NYC's law forcing disclosure was struck down in 2024 and affirmed in August 2026 (Daily Record).

Agent 4 — Inventory variance (inventory counting; portioning indirectly). Data: invoices, counts, recipes, item sales. It only works if someone keeps recipes and counts current, which is the work owners skip (operators: inventory software "wasted money" without accurate recipes and upkeep, PMQ). Build later, if at all.

Agent 5 — Forecast (forecasting). Toast IQ Team and Square Shifts already bundle forecasts; Lineup.ai sells them at $79 a location (vendor page, Lineup). Low priority.

Agent 6 — SOP nudges (training, SOP rollout). Two-minute SMS quizzes built from the owner's own procedures. Cheap, but incumbents fail here on "no time," not content: 62% of operators give 1–2 hours of training a week (vendor survey, 2026, Restaurant365).

Agent 7 — Hiring scheduler (new staff). Text replies and interview booking within minutes, because 36% of applicants drop out from slow replies (2025, Fountain/BBI); Chipotle's agent cut time-to-hire from 12 days to 4 (2026, Fortune). Do not let it score candidates; that triggers NYC's hiring bias-audit law (Local Law 144).

Problems software cannot solve without cameras: plate-by-plate portion size, drinks poured and never rung up (no POS trace), and phone use on shift. Even with cameras the evidence is bad: 18.8% portion error (2021), and Starbucks retired a camera-counting inventory agent after nine months because it "got less accurate over time" (2026, Fortune). The only vendor that monitors "phone use" uses human remote auditors, not AI (DTiQ).

Plumbing, in plain words. Start on Square: a personal access token is "suitable for production use" for one account with no approval (Square docs), and Square hosts an MCP server in beta (MCP = an open standard that lets an AI agent plug into other software, Square MCP). Toast costs the restaurant **$50 a month** for the portal integrations require, has no webhooks (webhook = a push notice when data changes) for orders, and its self-serve API price is unpublished (third-party help page, 2026,

MarginEdge). Clover may require app approval and a 30% revenue share (Clover docs). Deliver by text and email, not a dashboard: median software products see only 37% of users return after week one (2017, Mixpanel).

3. What it costs to run per restaurant

All figures are October 2026 list prices and my arithmetic. They exclude your time, support, sales, and any camera.

Design	Monthly cost	Arithmetic (one line)	Sources
Lean (Square, Haiku-class model, SMS digest, shared hosting)	**~$20–25**	AI tokens ~$1 (8.1M in + 0.56M out at $0.10/$0.50 per million) + SMS ~$15 (1,000 segments + number + 10DLC campaign fee) + hosting ~$5 + Square API $0	Anthropic, Twilio, Render
Typical (Toast, Sonnet-class model, 1,200 voice minutes)	**~$275**	AI ~$22 + Toast portal $50 + voice ~$179 (1,200 min at ~$0.149/min) + SMS ~$18 + hosting ~$5	Retell, MarginEdge on Toast
Heavy (multi-step agent loops, 3,000 voice minutes, own server)	**~$670**	AI ~$110 (5x tokens) + voice ~$450 + hosting $33 + Toast $50 + SMS ~$25	same as above, Vapi

The model bill is the smallest line; voice minutes and the POS vendor's gate fee dominate, and Toast's $50 portal charge alone exceeds the whole lean stack. Owners already pay $30–$150 a month for scheduling tools (2026, Homebase; Square), so a lean product at $79–99 leaves room to live. One licensing note: self-hosted n8n (a workflow tool) may be used "only for your own internal business purposes" (n8n license); running it for paying clients is a grey zone, so write automations in code or pay for n8n Cloud per client.

4. Solutions that already exist

Problem	Main tools	Price (2026)	Why owners use it	Why they don't
Voids, comps, cash theft	Toast/Square built-in reports; Solink, DTiQ (video + POS)	Reports free; Solink ~$175/location, DTiQ quote (GetApp)	Owners already track void % by hand (PMQ)	Reports have no alerts; Square cannot show badge vs code (Square Community); video tools take 2 months to set up (G2)
Time theft, scheduling	Homebase, Square Shifts, 7shifts	$30; $49; 7shifts Pro $89.99 (Homebase; G2)	Staff install the app that carries the schedule	Tier gating and repricing; 7shifts changed prices twice in Sep–Oct 2026 (PricingSaaS)
Inventory	MarketMan, MarginEdge, Square+MarketMan	$249; $350; $99 (MarketMan; MarginEdge)	Invoice scanning, not counting	"Never integrated with our POS"; cancellation traps (Capterra, 2025)

Problem	Main tools	Price (2026)	Why owners use it	Why they don't
Customer data, marketing	Square Plus loyalty, Owner.com, Toast IQ Grow	$49; $249–$499; $699 (Owner; NRN)	Dollar proof in 60–90 days	"Nickel and dime you for everything you want to connect" (Capterra, 2025)
Forecasting	Toast IQ Team, Lineup.ai	$60 + $9/employee; $79–$149 (Lineup)	Bundled in the POS	Sling's own docs say 30–40% error is "normal" (Sling)
Training, SOPs	Opus, Jolt, Trainual	Quote; $1,000 setup fee (Capterra)	Checklists on phones	"Almost $1k later… ZERO to show for it"; no time to complete
Hiring	Indeed, 7shifts Hiring, Homebase Hiring	Free; free; from $99 (Homebase)	Free covers one unit	Speed, not tools, is the gap
Phone answering	Yelp Host, Popmenu, Slang.ai	$149; $149; $399+ (Business Wire)	Captures missed reservations	Mishears accents and menu names; callers hang up (CloudTalk, 2026, competitor)
Portioning, phone use	Winnow, Leanpath, DTiQ SmartAudit	Quote-only, sold to chains (Winnow)	Hotel-scale food spend	No independent customer named; needs a scale, not just a camera

The 2025–2026 reviews repeat three complaints in every category: add-on fees, blocked cancellation, and support that disappears after the sale. Nobody in the dated set bought a tool because of AI (Capterra, 2025).

5. What people built that failed and why

Verified items only.

What	Date	What happened	Why
Staffjoy (restaurant scheduling)	Feb 2017	Shut down	Outbound email converted at ~0.1%; "not solving a core burning problem" (HN)
Leanpath at Lupa, NYC (waste scales)	2012	Chef ended pilot after ~2 months	Staff stopped weighing after the first week; it "took up valuable time" (WGBH)
Presto Automation (tablets, then voice AI)	Aug–Sep 2024	Delisted; sold by lenders, shareholders got nothing	54–70% of "AI" orders needed offshore human help (Restaurant Dive)
McDonald's / IBM drive-thru AI	June 2024	Ended at ~100 restaurants	Order accuracy in the low 80s vs 95% target (NRN)
DoorDash AI phone ordering	April 2025	Pilot ended	Not stated; customers switched vendors (Bloomberg)
Taco Bell voice AI (~500 units)	Aug 2025	Rollout slowed	Complaints and trolling; later moved to a new vendor (Restaurant Dive)
Starbucks "Automatic Counting" (camera inventory)	May 2026	Retired after ~9 months	Miscounted; "got less accurate over time" (Fortune)

What	Date	What happened	Why
Burger King voice AI (50–70 stores)	Sep 2026	Scaled back	"Drive offs because guests don't want to talk to a bot" (NRN)
Pizza Hut franchisee vs Dragontail AI	May 2026	$100M+ lawsuit (allegations only)	Mandated AI dispatch allegedly made drivers wait 15 minutes; managers lost control (The Register)

The pattern: AI that faces guests or takes control of an operation gets pulled when accuracy falls short. No owner-only, back-office analytics agent has a documented pullback yet. SpotOn's Profit AI (beta, June 2026) explicitly "does not make operational changes automatically" (vendor release, Business Wire). Recommend; do not act.

6. How to make sure owners like it and it does not fail

1. **Show facts, never a "thief" label.** In New York, "we fired her for stealing" is defamation per se if false (2025, Carey & Associates), and a 2006 court let a theft-accusation claim proceed where the facts suggested pretext (Kornsweig). Your output is "login X voided $412 across 9 checks between 10:40 and 11:55 pm, approved by code Y," sent to the owner only. 55% of Americans oppose AI in firing decisions, but only 39% oppose AI evaluating performance (2023, Pew). Call it "comp budget and cash-handling health."
2. **Month-to-month, cancel in one click.** New York's auto-renewal statute GOL §5-903 does not apply to renewal periods of one month or less (NY Senate), and GBL §527-a covers only personal purchases (NY Senate). Monthly billing sidesteps both and answers the top 2025 complaint: "trying to trap customers into paying longer than they want to" (Capterra, 2025).
3. **Write a data-export clause.** The standard SaaS contract (Common Paper CSA v2.1, 2026) gives the customer its data and deletion within 60 days, but has **no export right** (Common Paper). Add: 30-day CSV export of all flagged events after termination, the same if you shut down, and no AI training on their data.
4. **Set the pilot pass line in dollars before day one.** Measure voids + comps + discounts as a percent of net sales (sales after refunds) for 30 days before and 60 days after. Pass = a drop of at least 1 point of net sales, which on $1M of revenue is $10,000 a year against a $1,200 price. Fail = under 0.5 point, or the owner stops opening the digest. Keep it short: 57% of successful software buyers decided in 3 months or less; 54% of regretful ones took 5 months or more (2024, n=3,500, Capterra).
5. **Use design-partner terms, even with your friend.** (A design partner = an early customer who pays something and gives weekly feedback.) Charge 10–20% of the intended annual price up front, fix a term of about three months, hold a weekly 30-minute check-in, and end with "go paid or don't" (2025, First Round / Sierra). "Indefinite free access is an advisory relationship, not a validation signal" (2026, Bessemer). Use a written Design Partner Agreement so IP is clear (Common Paper).
6. **Give staff the New York notice.** Civil Rights Law §52-c requires written notice at hire, a signed acknowledgment and a posted sign before electronic monitoring, with penalties of $500, $1,000 and $3,000 per offense (NY Senate). Counsel alerts do not say whether reading POS logs counts (2022,

Kelley Drye); the notice costs one page, so include it. Tell staff nobody is punished for data showing up; that promise is the only documented staff-adoption tactic in food service (Leanpath, Notre Dame).

7. **No cameras in version one.** Besides the 18.8% error, New York bans video in restrooms and changing rooms (Labor Law §203-c), makes recording conversations without consent a felony (Penal Law §250.05), and bars mandatory fingerprint clocks (Labor Law §201-a) (NY Senate). NYC requires signage for any biometric capture with a private right of action (intro.nyc, LL 2021/3), and a state biometric bill passed the Senate 41–20 on June 3, 2026 (S1422-A).
8. **One narrow wedge first.** In B2B categories only about 5 brands reach 80% appearance in AI answers, and they are the incumbents (2026, Search Engine Land). You cannot out-bundle Toast. Own "void and comp monitoring for independents," prove it, then add Agent 2.
9. **Referral mechanics with real numbers.** Toast pays $1,000–$1,500 per referral, but only after the restaurant is live for 7 days, and "verbal referrals will not be accepted" (terms from March 1, 2026, Toast); 7shifts pays $100 after 60 days on a paid plan (7shifts); Homebase pays $100 to each side after 60 days (Homebase). Even Toast got only 20% of new locations from referrals (2022, SaaStr). So: a free month to both sides after the referred owner pays for 60 days, count a referral only then, and get the friend's name on a one-paragraph result. The NYC Hospitality Alliance costs $365 a year for one unit and runs owner roundtables and a vendor helpline (NYC Hospitality Alliance).
10. **Never do these.** Never bill a percent of "savings" you measure yourself; Toast's $0.99 fee lasted about ten days before the CEO called it "the wrong decision" (July 19, 2023, Toast). Never add hidden per-feature fees. Never send alerts to a staff group chat. Never let the AI be the only basis for a firing. Never make AI the only way a guest can order. Never pitch "AI" as the reason to buy: Popmenu's CEO said in May 2026 that "being first is not always the best" and that token costs for some tools rival the salaries they replace (Restaurant Dive). Owners' words are "catches comps and voids you'd miss."

7. Getting recommended by AI search tools: the honest truth

What the platforms say. Google: "There are no additional requirements to appear in AI Overviews or AI Mode," no special schema, inclusion "isn't guaranteed," llms.txt "neither help[s] nor harm[s]," and no third-party tool sees Google's systems (docs updated Dec 2025 and July 2026, Google; Google guide). OpenAI: "Any public website can appear in ChatGPT search" (OpenAI FAQ), and "Ads do not influence the answers ChatGPT gives you" (January 2026, OpenAI). Perplexity is shutting its ad program down through 2026 because ads make users "suspicious of everything" (February 2026, Campaign). Anthropic's web search docs describe always-on citations and no paid placement (Anthropic). Meta AI used Bing at launch (2023, Search Engine Land); its 2026 setup is not documented. Copilot is where Bing rank matters most (Ahrefs, 2025). No platform sells a mention inside the organic answer.

What studies show. Google AI Overview citations come from Google's top 10 results 76% of the time (2025) or 38% (2026), while ChatGPT citations overlap Google's top 10 only 6–8% (Ahrefs 2025; SEJ 2026). Across 75,000 brands, the strongest predictor of being mentioned in AI Overviews was how often the brand is named across the web (correlation 0.664), far ahead of backlinks (0.218) (Ahrefs, 2025). The one peer-reviewed study (KDD 2024) found that adding real statistics, quotations and source citations to a page raised its share of the answer by about 30–40%, keyword stuffing hurt, and lower-ranked pages gained most (arXiv). Review counts barely matter: 10% more G2 reviews ≈ 2% more citations (2025, G2).

What measurably helps, in order: rank in ordinary Google and Bing for one narrow query ("restaurant void and comp monitoring"); get named by third parties (trade press, and listicles like FitSmallBusiness, Upmenu and Coast, which dominated live results on October 16, 2026, FitSmallBusiness); list on G2/ Capterra/GetApp (cheap, sometimes cited); publish original pilot data with sources and quotes; show up as a named real person on Reddit and LinkedIn.

What is folklore: llms.txt, special "AI schema," content chunking, AI-rewritten copy, "rank #1 in ChatGPT" trackers, and Reddit seeding, which backfires publicly (Rippling was called out by subreddit moderators in March 2026, Foundation). Wikipedia is closed to new products without sustained independent coverage (Beutler Ink).

Why "#1" cannot be guaranteed. ChatGPT and Google returned the same brand list for the same prompt less than 1% of the time across 2,961 prompts (2026, SparkToro). In 100 runs of one B2B prompt, about 44 different brands appeared and only about 5 appeared in 80% of answers (2026, Search Engine Land). The answer is sampled fresh each time. Measure appearance rate over at least 50 runs per prompt per engine, and expect months, not weeks.

Gamma makes slides, documents and web pages from a prompt (Gamma). It does not search the web or recommend software. The correct targets are Google AI Overviews/AI Mode, ChatGPT, Perplexity, Claude, Gemini, Copilot and Meta AI.

---

8. What changed or got corrected since round one

- **Two judges, two cases.** Judge Gregory Woods let the delivery apps' suit against NYC's fee cap proceed (Restaurant Business); Judge Analisa Torres struck down the separate Customer Data Law.
- **Fee cap fight is over.** Int 0762-2024-B became **Local Law 79 of 2025** on May 31, 2025: caps stay at 15%/3%/5%, plus an optional "enhanced services" fee up to 20%, so the total can reach 43% (NYC Legistar). The apps dropped their suits in exchange (April 2025, Crain's).
- **7shifts repriced twice in Sep–Oct 2026.** Pro is $89.99 per G2 (October 2026); Premium swung between $134.99 and $179.99 (PricingSaaS; G2). The $69.99 figure was stale.
- **Toast Q2 2026 date** is **August 4, 2026**, not August 11 (SEC). Q3 2026 results are not yet released.
- **Chipotle settlement verified.** Up to about $20 million to about 13,000 NYC workers plus a $1 million penalty under the Fair Workweek Law (August 9, 2022, NYC Mayor's Office).
- **Toast $0.99 fee reversal dated:** July 19, 2023 (Toast).
- **NYC industry size updated.** About 280,000 restaurant and bar jobs in 2024 and "more than 23,000" establishments, versus 317,800 and 23,650 in 2019 (July 2026, OSC).
- **ACFE food-service row verified:** 53 cases, median loss $100,000 (2026, ACFE PDF).
- **New data-access facts:** Square's API cannot return voided items; Toast charges $50 a month for the integration portal; Square hosts an MCP server. **ServSafe Food Handler** costs $15 (ServSafe).

9. What we still could not verify

Not confirmed: Sling Pro's price; Black Box Intelligence's 2025 turnover from a primary page (a third-party roundup gives ~110% limited-service hourly); Toast's self-serve API price and partner fees; whether a Clover merchant token may run in production without app approval; whether Square exposes comp data through any API; whether §52-c covers reading POS logs; whether GOL §5-903 reaches SaaS at all; Meta AI's 2026 search provider; Toast's churn (no public restaurant-tech company discloses it); and any Reddit owner testimony, because Reddit was unreachable in both rounds. No survey gives an independent owner's willingness to pay in dollars. No study isolates restaurant-software queries in AI answers, and no independently audited case of a small software company gaining AI citations exists. The fates of Foodspace, Spotluck, Hang, Hyper, StockCount and Cheetah remain open. "None identified" for shutdowns of restaurant training or hiring startups in 2024–2026 means none found, not none occurred.

10. Bottom line

Do not build the all-in-one. Build a void, comp, discount and refund monitor that reads the POS, baselines each login and shift, and texts the owner a short ranked digest with time-stamped evidence. Start on Square because it needs no approval and no fee; treat Toast as phase two because it costs the restaurant $50 a month before you charge anything. Price it under $100 a month, month-to-month, with a data-export clause, and charge your friend something for a fixed 90-day term with the pass line written down first: at least 1 point of net sales recovered, or the pilot failed. Do not ship cameras, face recognition, audio, or a "thief" label in version one; New York law and the accuracy record both say no. Count a referral only when a second owner has paid for 60 days, because that is the standard every incumbent uses. For AI search tools, the only honest plan is to rank for one narrow phrase, get named by third parties, publish real pilot data with sources, and measure appearance rate over many runs; nobody can promise "#1," and Gamma was never on the list.